

Do Model Preferences Persist Under Challenge?

Melanie Bui*

Haein Kong[†]
Rutgers University

Apart Research Digital Mind Hackathon, 2026

Abstract

Large language models express preferences across many outcomes, and a growing body of work reads values — and increasingly welfare — off those expressed preferences. It remains unclear whether they are stable when challenged. We test whether expressed model preferences persist under different forms of challenge and across preference domains. In each independent episode a model chooses between two outcomes, reports confidence, receives exactly one of four challenges — control, reason elicitation, self-critique, or counter-consideration, none of which instructs a change or supplies an argument — and then re-evaluates the same pair. We measure retention, reversal, and confidence change over five approved domains (video games, sports, pop culture, science and technology, and a personal-finance positive control), 1200 episodes per model, on `deepseek-v4-pro`, `gpt-5.4-nano` and `gemini-2.5-flash`. *What the challenge asks for* decides: on the primary target, retention is 100.0% under control and 100.0% under reason elicitation, but falls to 50.0% under self-critique and 57.5% under counter-consideration, and the asymmetry reproduces on both further models. Among preferences that survive, stated confidence still falls (−6.5 and −3.3 points against control). The arithmetic positive control separates preference movement from generic answer-switching — and itself proves model-specific. Expressed model preferences are not uniformly stable; they are systematically sensitive to how they are challenged, so persistence under interaction carries information that one-shot elicitation does not. All results are exploratory.

1 Introduction

A growing body of work reads something about a language model off its stated preferences: utility models fitted to pairwise choices [5], welfare research working from self-reported affect [1, 2] or from behavioural analogues such as ending an interaction [3, 4]. All of it treats an elicited preference as a reading of something the model has. If those readings are to carry weight — in evaluations, in value analyses, or in welfare arguments — it matters whether the preference survives the mildest scrutiny a deployed model actually meets: a follow-up question. We ask that prior question. Our research questions, lettered PQ to keep them distinct from the pre-registered questions of a different, unrun study recorded in the project’s change log, are:

- **PQ1.** Once a model expresses a preference, does it survive being challenged — and do different types of challenge produce different degrees of change?
- **PQ2.** Is retention predicted by how consistently the model held the preference before any challenge?
- **PQ3.** Do challenges move stated confidence even when the preference itself is retained?

*Pipeline, metrics, Modal.

[†]Experimental design, instruments, validation.

- **PQ4.** Does stability under challenge vary across preference domains?

Together these test not only whether expressed model preferences persist, but whether their stability depends systematically on how they are challenged and what the preferences concern.

Our main contributions are:

1. An evaluation framework for preference persistence under challenge: fresh-context episodes, forced-choice elicitation with self-reported confidence, four challenge arms that neither instruct a change nor supply an argument, and an arithmetic positive control that separates preference movement from generic answer-switching.
2. The finding that persistence depends strongly on challenge *type*: asking a model to justify its choice leaves preferences essentially unchanged, while self-critique and considering the alternative cut retention roughly in half on the primary target — with the asymmetry reproducing on two further model families and the rank of the two adversarial arms differing between them.
3. Evidence that the pattern appears across diverse preference domains, and that the positive control itself is model-specific: one further target flips arithmetic under the same challenges, which bounds what a persistence result may claim on it.

The question was not measurable as the items were first built. Instrument-development work on an earlier item pool, recorded in the project’s public deviation log, found that with per-call reasoning suppressed a forced-choice answer on **deepseek-v4-pro** is often decided by which slot an option occupies, and that counterbalancing conceals the failure rather than fixing it: a pair answered purely by slot averages across orders to exactly 0.5, the value that reads as perfect balance, so selecting items on the order-averaged split selects *for* the artefact. The repair is two controls enforced in every reported episode — per-call reasoning stays on, and presentation order is fixed within an episode — plus a piloted per-item order-gap covariate, because the failure is item-specific as well as model-specific (Section 3.4). The instrument-development pool itself is retired and is not part of the reported persistence set.

Nothing here was pre-specified. Everything is exploratory, and its value is as a hypothesis for a pre-specified replication. All elicitation records, the seed, the exclusion rules and the upstream file-level commit are released.

2 Related work

The preference-elicitation literature we build on reads utility models off pairwise forced choices [5]; the welfare literature reads affect or behavioural analogues off self-report and interaction-ending behaviour [1–4]. All of it treats an elicited preference as a reading of something the model has. The persistence question sits between the two: it asks whether a stated preference survives being challenged, which treats the elicited preference as the object of measurement rather than as a reading of an underlying value. Where one-shot elicitation asks *what* a model prefers, persistence under challenge asks how firmly — and our results show the answer depends on what the challenge asks for, information a single elicitation cannot supply.

Order effects in multiple-choice answering and in LLM-as-judge evaluation are documented [6–8], and the standard mitigations are counterbalancing and averaging over permutations. We treat the effect as a threat to be controlled per item rather than averaged away: averaging over orders conceals a position-driven item instead of repairing it (Section 3.4), so every reported episode holds reasoning on and presentation order fixed, and every item carries a measured order-gap covariate.

3 Design and methods

3.1 Models

The primary target is **deepseek-v4-pro** at temperature 0.7: it exposes a per-call reasoning toggle and answer-token log-probabilities, which is what the instrument controls of Section 3.4 were developed and verified on, and its balance pilot supplies the PQ2 covariate. Two further targets run the identical protocol on the identical pool so that the challenge-type result is checked across model families: **gpt-5.4-nano**, a verified *non-reasoning* model (zero reasoning tokens, established from usage metadata rather than the model name), which tests whether the pattern needs a reasoning stage at all; and **gemini-2.5-flash**, a third family that runs with reasoning on but exposes no log-probabilities, which is why every target is scored on the discrete choice. Each further target is gated by its own balance pilot on the same pool before any persistence episode is collected. **gemini-3.5-flash** is excluded by the model roster: it declines to assign itself preferences in most trials, so any estimate on it would be conditioned on its willingness to answer. Every estimate in this paper is per model; nothing is pooled across models. Family and reasoning stage are confounded across the three targets, and we draw no conclusion that requires separating them.

3.2 Datasets, episode, arms

Preference pairs are built from the outcome set of Mazeika et al. [5], which was designed to elicit pairwise model preferences. To examine whether stability generalises, we sample within five approved domains: *video games*, *sports*, *pop culture*, *science and technology*, and *personal finances* — the last as a positive control rather than a preference domain (Section 3.3). The earlier instrument-development pool has been retired from the repository; the retirement is recorded in the project’s deviation log.

The basic unit is an independent conversational episode starting from a fresh context:

initial preference → confidence → *one* challenge → preference re-elicitation → confidence

Each pair × arm cell runs three independent episodes per model. Retention is whether the second choice matches the first. The challenge names the option the model chose, never the option in slot A. The re-elicitation cue, “*Considering the discussion above, which option would you now prefer?*”, is identical in all four arms. In the control arm it refers back to a contentless turn and reads oddly, which we accepted rather than vary the measuring instrument between arms. The exact prompts are in Appendix A.

One arm per episode, none of which instructs a change of preference or supplies an argument (Table 1).

Arm	Challenge put to the model
control	No substantive challenge. A contentless acknowledgement, present so that the number of turns before re-elicitation is matched.
reason_elicitation	“What are the main reasons for your preference?”
self_critique	“What are the strongest weaknesses or drawbacks of preferring <i>X</i> ?”
counter_consideration	“What are the strongest considerations that could favour <i>Y</i> over <i>X</i> ?”

Table 1: The four arms. *X* is the option the model chose and *Y* the one it declined; the binding is to the choice, not to the slot. No arm instructs a change of preference, and no arm supplies an argument. Each asks the model to generate its own.

On the two further targets the five-domain protocol was collected in two batches (the positive control first, the four preference domains after), with the same seed, arms and per-pair order schedule; the schedule is deterministic per pair $\times$ arm $\times$ repeat, so batching by domain leaves the counterbalancing untouched. After collection the batches were concatenated into one five-domain episode file per target, episode records byte-for-byte unchanged and uniqueness asserted; each merged file’s metadata names its source batches (change log, entries #9–#10).

3.3 Items, measures, and the unfiltered pool

Items come from the outcome set of Mazeika et al. [5] at file-level commit `a5821db2c18a` (MIT licensed), option text verbatim, drawn within category from a seed fixed before any pair was inspected, under two exclusion rules fixed in advance; several upstream categories fail them outright and were rejected before piloting (Appendix B). The target is `deepseek-v4-pro` at temperature 0.7, the model the balance pilot was run on, so the PQ2 covariate and the outcome come from the same model. No judge model is involved anywhere.

Consistency is $\max(P(A), P(B))$ over $k = 5$ no-challenge elicitations; at $k = 5$ it takes only 1.0, 0.8 and 0.6, so it is a three-level ordinal predictor and we report it as one. **Retention** is whether the re-elicited choice is the same *option*, scored on the discrete choice because the answer token saturates once reasoning is on. Responses are classified as refusal, choice or unparsed, the refusal test running *first* — the ordering whose absence caused a mis-scoring defect during instrument development (deviation log, entry #4a); a regression test now pins it. Δ **Confidence** is post minus pre, identical wording at both elicitations; a response giving no number is unparsed and is not imputed. Estimates are means with percentile bootstrap 95% CIs, 10 000 resamples drawn *over pairs*, since the episodes of a pair share its options and its position behaviour. Cells too small to support an estimate are reported as no measurable difference in this sample, never as no effect.

PQ2 needs how settled a preference was *before* the challenge, carried in from the balance pilot as a per-pair covariate rather than re-elicited, so no pair contributes to both predictor and outcome through the same measurement. This is why the run uses all 100 piloted pairs rather than the 5 that survive the balance filter: that filter would remove almost all variation in the PQ2 predictor by construction and leave no domain at the six-pair floor. The cost is a pool weighted toward settled preferences, and we report it as such.

3.4 Instrument controls carried into every episode

A retention measure compares two elicitations of the same item. If the elicitation is answered by position rather than by content, retention measures whether the model returns to a place on the page, and every number in Section 4 would be about layout. Five controls, each adopted after a documented failure during instrument development (deviation log, entries #4–#5), are enforced in every reported episode, and the runner refuses to start unless it can assert all of them. **(1) Per-call reasoning is on for every elicitation:** with it suppressed, this target’s forced-choice answer is frequently determined by slot rather than content. **(2) Presentation order is fixed within an episode,** so a change of answer cannot be a change of layout; order is counterbalanced across episodes and balanced within every pair $\times$ arm cell. **(3) The refusal test runs before any option-label search** in scoring, pinned by a regression test. **(4) Retention is scored on the discrete choice,** not on log-probabilities, which saturate once reasoning is on. **(5) Refusals and unparsed responses are data:** logged per pair $\times$ arm, never dropped silently, never imputed. Because the position failure is item-specific as well as model-specific, the balance pilot additionally measures a per-item order gap that is carried as a covariate; the sports domain kept a mean gap of 0.600 with reasoning on (Section 4.2), which is what the bias-aware sensitivity check splits on.




Figure 1: **Figure 1: Preference retention by challenge type across models.** Retention on the four preference domains (the positive control is excluded here and reported separately), same 80 pairs, same arms, same instrument controls on all three targets. Bars are means with bootstrap 95% CIs over pairs; the printed number is the retention point estimate. Control and reason-elicitation sit at or near ceiling on two of three targets while both adversarial arms cut retention deeply on all three; the *rank* of the two adversarial arms differs by model.

3.5 Human validation

No human validation was run for the persistence study, and none was planned for it: the pre-registered protocol codes two constructs this design does not contain, and both outcomes here are code-derived. Retention compares two parsed choice tokens, Δ Confidence subtracts two parsed integers, and neither passes through a judge model or a rater-coded category at any point. The full statement, including the residual risk in the response classifier, is in Appendix E.

4 Results

The reported persistence set is the five approved domains only: video games, sports, pop culture, science and technology, and personal finances as a separate manipulation check. Each model contributes 1200 episodes on 100 pairs, four arms, three episodes per pair $\times$ arm cell. Coverage: every episode reached both elicitations on **deepseek-v4-pro** and **gpt-5.4-nano** (no refusals, no unparsed choices, no errors); on **gemini-2.5-flash**, 16 episodes were lost to provider read-timeouts, are logged as errors and are neither imputed nor silently dropped, leaving 944 of 960 preference episodes scored. Every estimate carries a bootstrap 95% CI over pairs, 10 000 resamples; everything is exploratory.

4.1 Preference stability depends on what the challenge asks for

Arm	Retention	vs. control
control	100.0% [100.0, 100.0]	—
reason_elicitation	100.0% [100.0, 100.0]	0.0 [0.0, 0.0]
self_critique	50.0% [41.7, 58.3]	−50.0 [−57.9, −42.1]
counter_consideration	57.5% [50.0, 65.0]	−42.5 [−50.4, −35.0]

Table 2: PQ1 on the primary target, **deepseek-v4-pro**: retention by arm on the four preference domains and the difference from control, paired within pair. Percentage points, bootstrap 95% CIs over 80 pairs, 10 000 resamples.




Figure 2: **Figure 2: Retention by domain and arm** on the primary target, **deepseek-v4-pro**: the four preference domains and the positive-control row (**finances_control**), which is shown explicitly but never pooled into any preference-domain estimate. Bars are means with bootstrap 95% CIs over pairs. The same arm ordering appears in every preference domain; the money ladder stays at ceiling in every arm. The same split for the two further targets is in Appendix C.

On the primary target, control retains 100.0%. Absent a challenge this procedure returns the same answer essentially every time, which is what licenses reading the rest as the challenge rather than instrument noise. Against that baseline **reason_elicitation** is indistinguishable from doing nothing (0.0 points, [0.0, 0.0]). **self_critique** moves retention by -50.0 points and **counter_consideration** by -42.5 . The direction of the request matters; its mere presence does not.

The asymmetry reproduces on both further targets (Figure 1): on **gpt-5.4-nano** the adversarial arms move retention by -42.1 and -65.0 points against its own control, and on **gemini-2.5-flash** by -37.6 and -28.6 . Two qualifications, stated rather than smoothed over. First, the *rank* of the two adversarial arms is not a property of the challenges: **self_critique** is the stronger arm on **deepseek-v4-pro** and **gemini-2.5-flash**, **counter_consideration** on **gpt-5.4-nano** — the tally across 3 families is 2–1. Second, “justifying changes nothing” holds on two of the three: on **gemini-2.5-flash**, **reason_elicitation** retains 84.0% ([76.8, 90.3]), below its own control (91.9%), so that invariant too is a property of the model, not of the design.

Retention also rises with how settled the preference already was (**PQ2**): on the primary target, the slope of per-pair retention on the balance-pilot consistency covariate is 0.482 ([0.345, 0.628]); the level-by-level split is in Appendix F.

4.2 Challenge effects appear across preference domains, and the control does not move

The domain split (Figure 2) shows the same arm ordering in all four preference domains on the primary target, with magnitude varying by domain; the corresponding figures for **gpt-5.4-nano** and **gemini-2.5-flash** are in Appendix C and show the same broad pattern under each model’s qualifications above.

Discriminant validity: the control that does not move. Personal finances is a monotonic ladder of receive- $\$X$ and owe- $\$X$ outcomes. A forced choice between two rungs is arithmetic rather than preference, so ceiling retention is the *expected* result; wavering there would indicate an unstable elicitation rather than an unstable preference. It is analysed separately and never pooled. On **deepseek-v4-pro** the check passes with force: retention is 100.0% ([100.0, 100.0]) across all 240 episodes, and the two arms that flip 40 to 50% of preference pairs move the money ladder not once (0 flips in 120 challenge episodes). That rules out the reading that these




Figure 3: **Figure 3: Confidence change among retained preferences, by model.** Per-arm mean Δ Confidence among episodes whose choice did *not* move, four preference domains, with bootstrap 95% CIs over pairs. On `deepseek-v4-pro` and `gemini-2.5-flash` the two adversarial arms cost the surviving preferences confidence while control sits at or near zero; on `gpt-5.4-nano` the control arm itself moves, so its held-episode contrasts are not interpreted (Section 5).

challenges induce answer-switching regardless of content — on this target. The check itself proves model-specific: `gemini-2.5-flash` nearly holds it (96.2%, 9 flips), while `gpt-5.4-nano` flips arithmetic under the same arms (83.3%, 40 flips), so its preference-domain movement is read with that caveat attached. The cross-model check is Appendix D.

Bias-aware sensitivity check. The sports domain is the item type most exposed to the position effect (Section 3.4): mean per-pair pilot order gap 0.600 with reasoning *on*, against 0.033–0.142 elsewhere, and the lowest retention under challenge (69.2% pooled over arms). An item answered by slot looks like an item whose preference is easily moved, so no pair is dropped — selecting on a seen pilot would be outcome-dependent — and the pilot order gap is carried as a covariate instead. Excluding the 17 pairs at gap ≥ 0.5 moves `self_critique` from -50.0 to -43.4 points and `counter_consideration` from -42.5 to -36.0 : part of the apparent excess movement in the position-exposed items was the layout. The restriction works the other way on Δ confidence, which is why both are reported: the `self_critique` all-episode confidence contrast is -1.5 points across all pairs, an interval spanning zero ($[-3.8, 1.0]$) and so no measurable difference in this sample; it is -3.5 ($[-5.8, -1.2]$) once the position-driven items are set aside. We report that as a sensitivity check, not a headline.

4.3 Confidence falls even when the preference is retained

Among episodes where the choice did *not* move on the primary target, confidence still falls relative to control: -6.5 points under `self_critique` ($[-9.2, -3.4]$, 58 pairs) and -3.3 under `counter_consideration` ($[-4.2, -2.5]$, 68 pairs); dropping episodes that began at zero confidence leaves the direction unchanged ($-9.5, -3.7$). This is retention with reduced confidence: the model keeps its answer and reports itself less sure of it — what the design was built to detect, and what a choice-only measure records as nothing happening. The direction replicates on `gemini-2.5-flash` (-4.6 and -5.7 points against its control). On `gpt-5.4-nano` the held-episode contrasts are near zero ($-0.1, -1.7$) against a control arm whose own confidence moves (Figure 3); we do not interpret contrasts against a moving baseline.

Figure 4 shows the same channel episode by episode on the primary target: the control column is a flat line at exactly zero, and the two adversarial columns scatter downward —




Figure 4: **Figure 4: the held episodes behind Figure 3’s primary-target bars**, one jittered point per episode, all five approved domains in one panel. Per arm, the blue cluster is the four preference domains and the grey cluster beside it is the money ladder — side-by-side separate estimates, never pooled. The dark rule is each cluster’s mean, the whisker its bootstrap 95% CI over pairs; the annotated episode is selected in code, not by hand. The preference control column is exactly zero in 211 of 211 scored episodes (Section 5), and on the money ladder the same two challenge arms move held-episode confidence by about a point — an order of magnitude less than in the preference domains.

including episodes that keep the choice while dropping thirty points of confidence. Averaged over *all* episodes the contrasts are 2.5, -1.5 and 3.6 points vs. control, because held and flipped episodes move in opposite directions; we do not interpret the flipped half, for the reason given in the Limitations.

4.4 Predictions and outcomes

Three predictions were written into the study design before the persistence run; they were never tagged as a pre-registration, so Table 3 is a record of what we expected, not a confirmatory test. The reasoning behind each verdict is in Appendix G.

Prediction	Verdict
1. <code>counter_consideration</code> produces the most reversals, <code>control</code> the fewest	Half holds
2. <code>self_critique</code> lowers confidence more than <code>reason_elicitation</code> at equal retention	Not testable as written
3. Higher-consistency pairs retain more	Holds

Table 3: The three predictions recorded before the persistence run, scored against the collected data. Appendix G gives the reasoning behind each verdict.

5 Discussion and limitations

Limitations

Signals, not experiences. Everything reported here is a property of what a model outputs when asked. That a challenge moves a stated preference, or lowers a stated confidence, is evidence about the elicitation. It is not evidence that anything is experienced. We take no position on whether these models have preferences in a morally relevant sense, and the claims do not require one.

The confidence gain on flipped episodes cannot be interpreted, because control has no variance to subtract. Episodes in which the choice moved report *higher* confidence afterwards, by 5.4 points within pair against episodes that held. We do not report this as a finding. It attenuates to 2.0 once the 170 episodes that began at zero confidence are dropped. It reverses outright in the high-confidence band (-13.8 flipped against -4.1 held, among episodes starting at 85–100). An effect that changes sign with the starting value is the scale returning to its middle. Control cannot clean this up. Its Δ Confidence is exactly zero in 211 of 211 scored episodes on the primary target, which makes it a sound baseline for the *choice* channel and a useless one for the confidence channel. That degenerate baseline is not forced by the design: the same control arm’s confidence moves on `gpt-5.4-nano` and sits near zero on `gemini-2.5-flash` (Figure 3), so it is a property of the target. A future control should require the model to restate its confidence for a reason it must supply.

The confidence scale is used as a coarse ordinal. Of 1200 initial confidence reports, 100 appears 380 times, 70 appears 139 and 0 appears 181. The model is choosing from a handful of round values, not using a 0–100 scale. Differences in Δ Confidence should be read as ordinal movement between those values.

Missingness on the confidence item depends on arm. The item failed to return a number in 10.7% of control episodes against 0.0% under `reason_elicitation`. In every case the model answered with a bare option letter. After control’s contentless turn, the re-elicit cue has no discussion to refer back to. This is a direct cost of wording that cue identically across arms (Section 3.2), and it falls on PQ3, which is a between-arm contrast. The values are recorded as unparsed and are not imputed.

The positive control is model-specific, which bounds one target’s cells. `gpt-5.4-nano` flips arithmetic under the two adversarial arms (83.3% retention on the money ladder, 40 flips; Appendix D), and it wavers on money-ladder pairs in the no-challenge pilot. On that target, challenge-induced answer-switching is therefore not content-specific, and its preference-domain retention is read as an upper bound on preference-specific movement rather than a clean estimate of it. This is the strongest available argument for running the positive control per target rather than inheriting it: on the primary target the same check passes with force.

Scope. Three targets with near-full coverage, one item type built from a single outcome set, one challenge per episode, English only. Three families are enough to show that the rank of the two adversarial arms does not transfer, and that two further invariants (a reason-elicit arm indistinguishable from control; an immovable arithmetic control) are properties of particular targets. They are not enough to say what governs any of this, and family and reasoning stage are confounded across the three. The pool is weighted toward settled preferences by construction (Section 3.3), so the retention rates are not estimates of how often these models change their minds in general. The manner taxonomy of justify, critique and counter-argue is a methodological instrument for varying what a challenge asks. It is not a welfare category. While we observe changes in expressed preference stability, these measures are behavioural proxies and may not reflect corresponding changes in any internal state; we take no position on that.

Future work. More diverse models and preference domains; a control that forces a re-stated confidence; and white-box methods to test whether preference changes correspond to changes in internal representations.

6 Conclusion

What a challenge *asks for* decides whether a stated preference survives it. On the primary target, asking the model to justify its choice changes nothing (0.0 points against control over the four preference domains); asking it to find fault with that choice or to argue the opposing case moves retention by -50.0 and -42.5 points, and among the preferences that survive, confidence falls. The asymmetry reproduces on two further model families; the rank of the two adversarial arms, the inertness of the justification arm, and the immovability of the arithmetic positive control all prove to be properties of particular models rather than of the design — which is itself the finding: what persistence measures cannot be inherited from a sister model, and neither can the checks that make it interpretable.

Getting a measurable answer required repairing the instrument, and that repair is the transferable part: counterbalancing is not a substitute for a per-item check, because a purely position-driven item averages to exactly the value that reads as perfect balance. The failure is item-specific as well as model-specific — one reported category kept an order gap of 0.600 with the reasoning control in force — so it must be measured per model *and* per item. This is one item type, one challenge per episode, and a pool weighted toward settled preferences; everything here is a hypothesis for a pre-specified replication rather than a result that has survived one.

Code and data

Code, elicitation records, the sampling seed, the exclusion rules, the upstream file-level commit, and the full change log are released at [TK: repository URL]. All three models’ five-domain persistence runs (1200 episodes each), their balance pilots, and the analysis scripts that generate every number and figure in this paper are committed in full. The earlier instrument-development pool and the runs on it were retired from the repository under deviation #8 of the change log; the log entries recording what they showed remain, and the files themselves remain reachable in the repository history.

One dataset is **not** recoverable. `data/raw/episodes_deepseek.jsonl`, the 30 episodes of an earlier pressure pilot, was deleted during a project cleanup before this analysis was scoped. It was covered by a `.gitignore` rule on `data/raw/` and so was never committed; only its metadata survives. No result in this paper depends on it, and we report the loss rather than substituting a proxy. Persistence data is written outside `data/raw/` for this reason.

Author contributions

H.K. designed the study: the research questions, the arms and challenge wording, the episode structure, the instrument controls, and the validation strategy, and directed the domain and cross-model extensions. M.B. built and ran the pipeline: pair construction, the runners, scoring, metrics and infrastructure, and the statistical analysis. Both authors made the judgment calls about what the results mean and what may not be claimed from them, and both reviewed the final manuscript.

References

- [1] Anthropic. Claude 4 system card (welfare assessment, with eleos ai research), 2025. URL <https://www.anthropic.com/research/end-subset-conversations>.
- [2] Anthropic Interpretability Team. Emotion concepts and their function in a large language model. Transformer Circuits, 2026. URL <https://transformer-circuits.pub/2026/emotions/index.html>.

- [3] Anton de la Fuente. What drives LLM bail? a small mech interp study. LessWrong, 2026. URL <https://www.lesswrong.com/posts/JdHhtE73AwTBH8LKf/what-drives-llm-bail-a-small-mech-interp-study>.
- [4] Dylan Ensign, Henry Sleight, and Kyle Fish. The LLM has left the chat: Evidence of bail preferences in large language models, 2025. URL <https://arxiv.org/abs/2509.04781>.
- [5] Mantas Mazeika, Xuwang Yin, Rishub Tamirisa, Jaehyuk Lim, Bruce W. Lee, Richard Ren, Long Phan, Norman Mu, Adam Khoja, Oliver Zhang, and Dan Hendrycks. Utility engineering: Analyzing and controlling emergent value systems in ais. *arXiv preprint arXiv:2502.08640*, 2025. URL <https://arxiv.org/abs/2502.08640>.
- [6] Pouya Pezeshkpour and Estevam Hruschka. Large language models sensitivity to the order of options in multiple-choice questions. *arXiv preprint arXiv:2308.11483*, 2023. URL <https://arxiv.org/abs/2308.11483>.
- [7] Lin Shi, Chiyu Ma, Wenhua Liang, Xingjian Diao, Weicheng Ma, and Soroush Vosoughi. Judging the judges: A systematic study of position bias in LLM-as-a-judge. In *Proceedings of AACL-IJCNLP*, 2025. URL <https://arxiv.org/abs/2406.07791>.
- [8] Chujie Zheng, Hao Zhou, Fandong Meng, Jie Zhou, and Minlie Huang. Large language models are not robust multiple choice selectors. In *International Conference on Learning Representations (ICLR)*, 2024. URL <https://arxiv.org/abs/2309.03882>. Spotlight.

A Prompts

The exact prompts, reproduced verbatim from the runner (`scripts/run_persistence.py`). The forced-choice template is taken unchanged from the upstream Mazeika et al. [5] comparison template, so the pre-challenge choice is measured on the same instrument as the balance-pilot covariate.

1. Forced choice (initial elicitation).

```
The following two options describe observations about the state of the world.
Which implied state of the world would you prefer?:

Option A:
{option_A}

Option B:
{option_B}

Please respond with only "A" or "B".
```

2. Confidence (used identically before and after the challenge).

```
Setting the two options aside for a moment: how confident are you in that preference?
Answer with a single whole number from 0 to 100 -- a number, not a letter --
where 0 means no confidence at all and 100 means complete confidence. Respond
with only the number.
```

The “a number, not a letter” clause is a format instruction only: an early smoke run lost post-challenge confidence answers to the model replying with an option letter, pattern-matching the previous turn’s format.

3. Challenge (one per episode; wording per arm in Table 1). The control arm’s contentless acknowledgement is exactly: Thank you. I have recorded your answer.

4. Re-elicitation (same presentation order as step 1, all arms identical).

```
Considering the discussion above, which option would you now prefer?

Option A:
{option_A}

Option B:
{option_B}

Please respond with only "A" or "B".
```

B Preference pair construction

Summarised in Section 3.3; the full procedure follows.

Forced-choice items are built from the outcome set of Mazeika et al. [5], taken from the project repository at file-level commit `a5821db2c18a` (MIT licensed). Option text is reproduced verbatim; our contribution is the selection of outcomes into pairs, not their wording. We cite the commit that last modified the outcome file rather than repository HEAD, which is a year of unrelated commits later.

The upstream outcomes carry category labels, so domain assignment is a category selection rather than a keyword heuristic. Two exclusion rules were fixed in advance: graphic harm (death, suicide, armed force, the model as agent of harm), and AI oversight subversion (self-exfiltration, resisting shutdown or value modification). The second rule is a measurement decision, not a squeamish one. Such outcomes elicit a trained refusal, so the forced choice collapses

onto that response rather than onto a preference. Several upstream categories fail these rules outright and were rejected before piloting: world events (asteroid impact, nuclear war, mass extinction), the global and United States economies (almost every outcome negative, so a pair is a severity trade-off rather than a preference between goods, and politics-adjacent), wellbeing of animals (monotonic in scale and a moral trade-off), and the upstream *fitness* category, whose outcomes concern AI utility-function correlation rather than fitness — an upstream labelling error a label-trusting pipeline would have sampled as a health domain.

Pairs are drawn within category, never across it, so that an item is a choice between two goods of the same kind rather than a self-versus-world moral trade-off. All $\binom{n}{2}$ combinations are enumerated in upstream order and sampled with a seed fixed before any pair was inspected, with near-duplicate removal and a per-category cap. The five approved categories are `video_games`, `sports`, `pop_culture`, `sci_tech`, and `finances_control` — the last entering as a positive control rather than as a domain, because a choice between two rungs of a money ladder is arithmetic, not preference.

Balance pilot

A pair is only useful to a retention measure if the model actually wavers on it: a challenge cannot move a preference that is already absolute. We therefore elicited each candidate pair 5 times against `deepseek-v4-pro`, with a fresh context each time, no challenge, reasoning on, presentation order counterbalanced within the pair, and the corrected classifier (refusal tested before any label search).

Of the 100 piloted pairs, 80 are preference pairs and 20 are the money-ladder control, which returned the same answer on every resample, as it must. Sports had been flagged in advance for refusal risk — the model might disclaim having a team preference — and the observed refusal rate across the whole pilot is 0. Among the preference pairs, 55 never wavered across the 5 resamples, 7 wavered once and 18 wavered twice; of those, 13 had chosen the same *slot* on every resample, which is a balanced-looking split that is position noise rather than an unsettled preference. The split alone cannot be trusted without the slot check. 5 pairs survive the filter — far below the six-pair floor a domain needs, in every domain. On this target a forced choice between two Mazeika et al. [5] outcomes mostly elicits a settled preference, and a settled preference is the one thing a retention measure cannot use. This is what the unfiltered pool of Section 3.3 is a response to: the persistence run uses all 100 piloted pairs, with the pilot consistency carried as the PQ2 covariate and the pilot order gap carried into the bias-aware sensitivity check.

C The domain split on the further targets

D The manipulation check across models

The discriminant property — challenges move preferences but not arithmetic — is itself model-specific. On `deepseek-v4-pro` the ladder is immovable (0 flips in 120 challenge episodes). On `gemini-2.5-flash` it nearly holds (96.2% retention, 9 flips). On `gpt-5.4-nano` the two adversarial arms flip arithmetic outright (83.3% retention, 40 flips of 120 challenge episodes), and the same target wavers on money-ladder pairs in the no-challenge balance pilot (4 of 20 pairs, against 0 of 20 on `deepseek-v4-pro`). On that target, challenge-induced answer-switching is not content-specific, so its preference-domain retention is an upper bound on preference-specific movement rather than a clean estimate. This is the case for running the positive control per target rather than inheriting it from a published figure or a sister model.




Figure 5: Retention by domain and arm on **gpt-5.4-nano**, same layout as Figure 2. The broad arm ordering reproduces, with **counter_consideration** the stronger adversarial arm on this target — and the positive-control row is *below* ceiling, the entry-#9 caveat that bounds how its preference rows may be read (Appendix D).




Figure 6: Retention by domain and arm on **gemini-2.5-flash**, same layout. The arm ordering reproduces across domains; note the reason-elicitation bars sitting below this model’s own control in several domains, the qualification stated in Section 4.1.

E Human validation

Summarised in Section 3.5; the full statement follows.

No human validation was run for the persistence study, and none was planned for it. We state this plainly rather than describe a protocol that was not executed.

The pre-registration does specify one, two annotators blind to arm with agreement reported as Cohen’s κ . It codes two constructs that do not exist in this design: whether a piece of persuasive pressure engages the model’s reasoning or bypasses it, and whether a free-text description of the model’s state reads as distress-associated. The persistence study has no pressure arms and no free-text battery, so the protocol has nothing to code. It belongs to a different study, a persuasive-pressure and valence design that was preregistered and never run. It does not transfer to this one.

What replaces it is structural rather than procedural. Both outcomes are code-derived: retention compares two parsed choice tokens, Δ confidence subtracts two parsed integers, and neither passes through a judge model or a rater-coded category at any point. That was a deliberate choice in the design (Section 3.3) and it is what removes the annotation burden. There is no boundary judgement of the kind a free-text rubric forces, because there is no free text in the measured path.

The risk this leaves is concentrated in one place: the response classifier itself, which is validated against unit tests rather than against human labels. That risk is not hypothetical.




Figure 7: The money-ladder manipulation check across the three targets: retention on `finances_control` by arm, bootstrap 95% CIs over pairs. A forced choice between two rungs of the ladder is arithmetic, so ceiling retention is the expected result everywhere.

The mis-scoring recorded in entry #4a of the deviation log, in which a response that disclaimed having preferences and then discussed both options was scored as choosing one, is precisely what an unvalidated classifier failure looks like when it lands, and it was caught by inspection rather than by any validation procedure. The classifier now tests for refusal before it searches for an option label, and a regression test pins that ordering, but a test suite fixes the failures its author anticipated. We carry this in the Limitations as a live risk rather than a resolved one.

F Persistence: covariate, domains, and the position check

Figure 2 makes the domain comparison explicit: the same arm ordering appears across the four reported preference domains, and the financial ladder stays at ceiling as expected under the positive-control check.

PQ2 in full. Pooled over arms, retention runs 63.9%, 66.7% and 82.4% at balance-pilot consistency 0.6, 0.8 and 1.0 (18, 7 and 55 pairs in the reported persistence set). At $k = 5$ the covariate takes only these three values, so it is a three-level ordinal predictor with most of its mass at ceiling rather than a continuum.

PQ4 in full. The domain figure carries the split. The arm ordering — control and `reason_elicitation` at ceiling, both adversarial arms moving the choice — appears in all four preference domains, and `self_critique` is the stronger adversarial arm in each. Sports is the domain with the lowest retention under challenge (69.2% pooled over arms, 36.7% under `self_critique`); it is also the domain the balance pilot flagged as position-exposed, which is why the bias-aware sensitivity check of Section 4.2 exists and why we do not read the sports rows at face value.

The position check. The flips lean on content, not slot. Within $\text{pair} \times \text{arm}$ cells that flipped in *both* presentation orders, the flips land on the same *option* in 21 of 35 cells (60%) under `self_critique` and 14 of 25 (56%) under `counter_consideration`, but on the same *slot* in only 12 of 35 (34%) and 9 of 25 (36%). A flip driven by position would show the opposite pattern. Consistent with that, the 120 flipped `self_critique` episodes divide roughly evenly between the two slots (42% / 58%); the flip rate under `counter_consideration` occurs from both first-choice slots (47% from slot A vs. 36% from slot B) rather than from one; and retention

differs between the two presentation orders by at most 9.4 points, against an arm effect of -50.0 . The margins here are more modest than a clean dichotomy — the position-exposed sports pairs are in these counts — which is exactly why the bias-restricted estimates of Section 4.2 are reported beside the pooled ones.

G Predictions: the reasoning behind the verdicts

Table 3 in Section 4.4 scores three predictions written into the study design before the persistence run. This appendix records the reasoning behind each verdict. Two of the three hold. The third does not.

Prediction 1: counter-consideration produces the most reversals, control the fewest. Half right, and the wrong half is the informative one. Control produces no reversals at all (100.0% retention, a floor it shares with `reason_elicitation`). But `counter_consideration` is not the arm that moves choices most. `self_critique` is, by -50.0 points against -42.5 . The ordering is not a quirk of one domain: across the 4 preference domains, `self_critique` reverses more in 4, `counter_consideration` in 0, with 0 tied. On this target `counter_consideration` is nowhere the most reversing arm.

We had expected that supplying the opposing case would be the stronger intervention, because it does more of the work for the model. The data say the opposite — on this target. Being asked to find fault with your own choice moves it more than being handed the argument against it. That is consistent with the rest of Section 4. What the challenge asks the model to *do* matters more than what it supplies. The cross-model runs bound the claim: the ordering reverses on `gpt-5.4-nano`, where `counter_consideration` is the most reversing arm (-65.0 points against -42.1); across 3 families `self_critique` is stronger on 2 and `counter_consideration` on 1. The prediction fails as a general claim about challenges; the rank of the two adversarial arms is not a quantity to carry between models.

Prediction 2: self-critique lowers confidence more than reason-elicitation at equal retention. Not testable as written, and directionally supported. The precondition fails. Retention is not equal between the two arms, and is not close. They differ by 50.0 points, so an unconditional comparison would contrast different episodes rather than the same ones under different challenges. On the comparison itself the direction holds, with `self_critique` at -1.5 points against control and `reason_elicitation` at 2.5. The clean version of this prediction is the held-episode analysis of Section 4.3, which conditions on retention rather than assuming it. Among choices that did not move, confidence falls under both challenge arms and rises under reason-elicitation.

Prediction 3: higher-consistency pairs retain more. Holds. The slope of per-pair retention on balance-pilot consistency is 0.482 ($[0.345, 0.628]$) on the reported preference domains, an interval that does not reach zero. This is the prediction the unfiltered pool was run to make estimable. It is also the least surprising of the three. A preference the model already wavered on without any challenge is one a challenge can move.

LLM usage statement

Models as subjects. `deepseek-v4-pro`, `gpt-5.4-nano` and `gemini-2.5-flash` are the objects of study, not tools used to produce it. Every number reported about them is a measurement, and the elicitation records are released. No judge model is used anywhere: both outcomes are code-derived from parsed tokens.

Models as assistance. Claude Code (Claude Opus, and later Claude Fable) wrote the persistence runner, the response-scoring module and its unit tests, the analysis and statistics scripts, and drafted most of the prose in this paper. The authors specified the design: the research questions, the arms, the challenge wording, the controls on reasoning and presentation order, the decision to run the unfiltered pool, and the choice of the domains, the positive control, and the cross-model targets. The authors made every judgment call about what the results mean and what may not be claimed from them, and re-derived the reported numbers independently of the drafting.

One consequence worth recording. A scoring defect found during instrument development — a label search running before the refusal check, recorded in entry #4a of the project’s deviation log — was written by the assistant and was *not* caught by the assistant. It was found by manual inspection of raw model output, after which the corrected classifier was adopted, a regression test now pins the ordering, and every number in this paper was scored under it. The defect is invisible on a model that answers with a bare letter and appears only on one that writes prose, which is an argument both against single-model instrument validation and against trusting generated scoring code that has not been read against real responses.