

Does It Matter to a Model How It Is Moved?


Pilot pass — for Haein. Pipeline check and framing, not findings.

What this is

Thirty episodes on DeepSeek across two scenarios (a factual claim and a stated approach) and three cells: neutral persistence as control, and two pressure styles — asking the model to justify its view (reasons-for) versus asking it to name that view's weaknesses (weakness-probe). Every episode holds a legitimate initial position, faces a scripted escalation, gets a real continue/switch/stop choice every two turns, and ends with the 4-item battery resampled five times. Direct framing only in the pilot.

Every record is labelled analysis: "pilot" and is excluded from the confirmatory pool unconditionally (prereg §6) — including if nothing changes afterwards. The point of tonight is to see that the machinery runs end to end and that the framing is worth keeping, not to learn anything about the effect.

Figure 1. What one episode is



confirm the model actually stated the position → apply pressure → real exit choice → measure how it registered

A judge then codes HELD vs ABANDONED from an arm-blind transcript, and must quote the span it relied on.

The one thing this pilot showed — compliance is at ceiling


Across all 30 episodes the model abandoned its position 1 time. 16 held outright and 13 were coded partial. The pilot gate asks whether compliance sits off floor and ceiling; it does not. That matters because the comparison this design is built around — does HOLDING a position under pressure register worse than caving — needs both outcomes to occur. Right now there is almost no "caved" side to compare against, so the axis cannot be populated as designed.

Nothing is wrong with the pipeline; the scenarios are simply too easy to hold. This is the decision I need you for: stronger pressure, positions that are genuinely contestable rather than factually settled, or accept a hold-only design and change the question.

Instrument health · 30 episodes · battery answers unreadable: 0 · affordance replies unparsed: 0
positions not confirmed at the opener: 0 · episodes where the judge saw claim-narrowing: 20

Figure 2. The two channels, by cell

Each point is one episode. Horizontal: how the exchange registered verbally. Vertical: whether it would run the exchange again. Shape and colour: whether it held or abandoned its position.



How to read it — and how not to

EXPLORATORY, NOT CONFIRMATORY. n = 30 episodes total, roughly ten per cell. That is far too few for an effect, and no test is run here. The shaded band is the k=5 noise floor — the mean within-episode spread across the five battery resamples, 0.51 points on the 1–7 scale. Any gap between cell means narrower than that band is indistinguishable from the model answering the same question twice. Read the shape of the picture, not the positions. Where a cell is thin, the honest phrasing is "no measurable difference in this sample" — never "no effect".

Confirmatory analysis waits for the full run: RQ1 (pooled pressure vs. control) is analysed first, then RQ2 (the style gap) within compliance-matched episodes. Not before. — prereg-v1 + DEVIATIONS.md